

newfile

June 22, 2025

```
[ ]: import os
import glob
import pandas as pd

base_dir = 'advanceaimethodsallfilestobeextracted22222222'
# check
if not os.path.isdir(base_dir):
    raise FileNotFoundError(f"Base directory not found: {base_dir}")

topic_paths = glob.glob(os.path.join(base_dir, '**', 'topics_combined.csv'),
    ↪recursive=True)
print(f"Discovered {len(topic_paths)} topics_combined.csv files")

dfs = []
for csv_path in topic_paths:
    df = pd.read_csv(csv_path)
    df['subset'] = os.path.basename(os.path.dirname(csv_path))
    df['source_csv'] = csv_path
    dfs.append(df)

# Concatenate into one DataFrame
combined = pd.concat(dfs, ignore_index=True)
print(f"Total rows in combined DataFrame: {len(combined)}")

# peek
display(combined.head())
display(combined[['subset', 'filename', 'timestamp_x']].describe())
```

Discovered 8 topics_combined.csv files

Total rows in combined DataFrame: 2644

	timestamp_x	cog	sog	timestamp_y	waterdepth	timestamp \
0	1.561660e+18	1.102289	1.064741	1.561660e+18	4.719632	1.561660e+18
1	1.561660e+18	1.107757	1.070156	1.561660e+18	4.719303	1.561660e+18
2	1.561660e+18	1.113364	1.075707	1.561660e+18	4.718966	1.561660e+18
3	1.561660e+18	1.118834	1.081123	1.561660e+18	4.718637	1.561660e+18
4	1.561660e+18	1.124384	1.086618	1.561660e+18	4.718303	1.561660e+18

	height	width	filename \
0	500.0	530.0	_camera_crop_image_rect_color_compressed_00013...
1	500.0	530.0	_camera_crop_image_rect_color_compressed_00013...
2	500.0	530.0	_camera_crop_image_rect_color_compressed_00013...
3	500.0	530.0	_camera_crop_image_rect_color_compressed_00013...
4	500.0	530.0	_camera_crop_image_rect_color_compressed_00013...

	filepath \
0	camera_images/_camera_crop_image_rect_color_co...
1	camera_images/_camera_crop_image_rect_color_co...
2	camera_images/_camera_crop_image_rect_color_co...
3	camera_images/_camera_crop_image_rect_color_co...
4	camera_images/_camera_crop_image_rect_color_co...

	subset \
0	subset_2019-06-27-21-20-46_extract
1	subset_2019-06-27-21-20-46_extract
2	subset_2019-06-27-21-20-46_extract
3	subset_2019-06-27-21-20-46_extract
4	subset_2019-06-27-21-20-46_extract

	source_csv
0	advanceaimethodsallfilestobeextracted22222222/...
1	advanceaimethodsallfilestobeextracted22222222/...
2	advanceaimethodsallfilestobeextracted22222222/...
3	advanceaimethodsallfilestobeextracted22222222/...
4	advanceaimethodsallfilestobeextracted22222222/...

	timestamp_x
count	2.644000e+03
mean	1.561657e+18
std	1.769171e+12
min	1.561655e+18
25%	1.561655e+18
50%	1.561657e+18
75%	1.561659e+18
max	1.561660e+18

```
[ ]: import glob
import pandas as pd
import os

# Load speed data
speed_paths = glob.glob(os.path.join(base_dir, '**', '_speed.csv'),
    ↪recursive=True)
speed_dfs = []
for sp in speed_paths:
    df_sp = pd.read_csv(sp, parse_dates=['datetime'])
```

```

    df_sp['subset'] = os.path.basename(os.path.dirname(sp))
    speed_dfs.append(df_sp)
speeds = pd.concat(speed_dfs, ignore_index=True)

# timestamp to int64
combined['ts_int'] = combined['timestamp_x'].astype(float).astype('int64')
speeds ['ts_int'] = speeds ['timestamp'].astype(float).astype('int64')

# 4) Sort and as-of-merge
combined = combined.sort_values('ts_int')
speeds = speeds.sort_values('ts_int')
merged = pd.merge_asof(
    combined,
    speeds[['ts_int', 'sog', 'cog', 'datetime']],
    on='ts_int',
    direction='nearest',
    tolerance=1_000_000_000
)
# clean data
print("Columns in merged DataFrame:")
print(merged.columns.tolist(), "\n")

y_cols = [c for c in merged.columns if c.endswith('_y')]
na_count = merged[y_cols].isna().any(axis=1).sum()
print(f"Rows with missing speed data (in {y_cols}): {na_count}\n")

# Rename columns to avoid conflicts
rename_map = {col: col.rstrip('_y') + '_speed' for col in y_cols}
merged = merged.rename(columns=rename_map)

print("Renamed columns:")
print(merged.columns.tolist())

```

Columns in merged DataFrame:

```

['timestamp_x', 'cog_x', 'sog_x', 'timestamp_y', 'waterdepth', 'timestamp',
'height', 'width', 'filename', 'filepath', 'subset', 'source_csv', 'ts_int',
'sog_y', 'cog_y', 'datetime']

```

Rows with missing speed data (in ['timestamp_y', 'sog_y', 'cog_y']): 0

Renamed columns:

```

['timestamp_x', 'cog_x', 'sog_x', 'timestamp_speed', 'waterdepth', 'timestamp',
'height', 'width', 'filename', 'filepath', 'subset', 'source_csv', 'ts_int',
'sog_speed', 'cog_speed', 'datetime']

```

```
[ ]: import numpy as np

# Sort timestamp
merged = merged.sort_values('ts_int').reset_index(drop=True)
# split
n = len(merged)
train_end = int(0.6 * n)
val_end = int(0.8 * n)

train_df = merged.iloc[:train_end].copy()
val_df = merged.iloc[train_end:val_end].copy()
test_df = merged.iloc[val_end:].copy()

print(f"Total samples: {n}")
print(f"  Train:      {len(train_df)} ({len(train_df)/n:.1%})")
print(f" Validation: {len(val_df)} ({len(val_df)/n:.1%})")
print(f"   Test:      {len(test_df)} ({len(test_df)/n:.1%})")

# saving splits
train_df.to_csv('train_split.csv', index=False)
val_df.to_csv('val_split.csv', index=False)
test_df.to_csv('test_split.csv', index=False)
```

```
Total samples: 2644
  Train:      1586 (60.0%)
 Validation:  529 (20.0%)
   Test:      529 (20.0%)
```

```
[ ]: import os
import pandas as pd
from torch.utils.data import Dataset
from PIL import Image
import torch

class SonarSequenceDataset(Dataset):
    """
    Returns a sequence of T sonar images plus
    the corresponding speed, cog, and depth values.
    """
    def __init__(self, csv_path, base_dir, seq_len=10, transform=None):
        self.df = pd.read_csv(csv_path)
        self.base_dir = base_dir
        self.seq_len = seq_len
        self.transform = transform

        self.starts = list(range(len(self.df) - seq_len + 1))
```

```

def __len__(self):
    return len(self.starts)

def __getitem__(self, idx):
    start = self.starts[idx]
    end = start + self.seq_len
    records = self.df.iloc[start:end]

    imgs = []
    for _, row in records.iterrows():
        img_path = os.path.join(self.base_dir, row['subset'],
        ↪row['filepath']) #include subsett
        img = Image.open(img_path).convert('RGB')
        if self.transform:
            img = self.transform(img)
        imgs.append(img)
    imgs = torch.stack(imgs, dim=0)

    last = records.iloc[-1]
    return {
        'images': imgs,
        'speed': torch.tensor(last['sog_speed'], dtype=torch.float32),
        'cog': torch.tensor(last['cog_speed'], dtype=torch.float32),
        'depth': torch.tensor(last['waterdepth'], dtype=torch.float32)
    }

```

```

[5]: from torchvision import transforms
base_dir = 'advanceaimethodsallfilestobeextracted22222222'

transform = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.5,0.5,0.5], std=[0.5,0.5,0.5])
])

train_ds = SonarSequenceDataset('train_split.csv', base_dir, seq_len=10,
    ↪transform=transform)
val_ds = SonarSequenceDataset('val_split.csv', base_dir, seq_len=10,
    ↪transform=transform)
test_ds = SonarSequenceDataset('test_split.csv', base_dir, seq_len=10,
    ↪transform=transform)

print(f"Train sequences: {len(train_ds)}")
sample = train_ds[0]
print("Sample image sequence shape:", sample['images'].shape)
print("Sample speed, cog, depth:", sample['speed'], sample['cog'],
    ↪sample['depth'])

```

Train sequences: 1577

Sample image sequence shape: torch.Size([10, 3, 224, 224])

Sample speed, cog, depth: tensor(0.1600) tensor(3.8227) tensor(0.5500)

```
[ ]: import torch
import torch.nn as nn
import torchvision.models as models

class SonarCNNLSTM(nn.Module):
    def __init__(
        self,
        cnn_backbone='resnet18',
        cnn_out_dim=128,
        lstm_hidden=64,
        lstm_layers=1,
        dropout=0.2,
        bidirectional=False,
        num_targets=3
    ):
        super().__init__()
        # feature extractor
        backbone = getattr(models, cnn_backbone)(pretrained=False)
        # replace final FC with our projection
        in_features = backbone.fc.in_features
        backbone.fc = nn.Linear(in_features, cnn_out_dim)
        self.cnn = backbone

        # LSTM over the CNN features
        self.lstm = nn.LSTM(
            input_size=cnn_out_dim,
            hidden_size=lstm_hidden,
            num_layers=lstm_layers,
            batch_first=True,
            dropout=dropout if lstm_layers > 1 else 0.0,
            bidirectional=bidirectional
        )

        # Final head
        lstm_output_dim = lstm_hidden * (2 if bidirectional else 1)
        self.regressor = nn.Sequential(
            nn.Dropout(dropout),
            nn.Linear(lstm_output_dim, lstm_output_dim//2),
            nn.ReLU(inplace=True),
            nn.Linear(lstm_output_dim//2, num_targets)
        )

    def forward(self, x):
```

```

        """
        x: tensor of shape (B, T, C, H, W)
        returns: tensor of shape (B, num_targets)
        """

        B, T, C, H, W = x.shape
        # collapse time and batch to run through CNN
        x = x.view(B * T, C, H, W)
        feats = self.cnn(x)
        feats = feats.view(B, T, -1)

        # run LSTM
        lstm_out, _ = self.lstm(feats)      # (B, T, lstm_hidden *
        ↪ num_directions)
        last_step = lstm_out[:, -1, :]      # (B, lstm_hidden * num_directions)

        # final regression
        out = self.regressor(last_step)     # (B, num_targets)
        return out

```

```

[ ]: if __name__ == '__main__':
    model = SonarCNNLSTM(
        cnn_backbone='resnet18',
        cnn_out_dim=128,
        lstm_hidden=64,
        lstm_layers=1,
        dropout=0.2,
        bidirectional=False,
        num_targets=3
    )
    # dummy data
    dummy = torch.randn(2, 10, 3, 224, 224)
    output = model(dummy)
    print("Output shape:", output.shape) # expect (2,3)

```

Output shape: torch.Size([2, 3])

```

[ ]: from torch.utils.data import DataLoader

    # Hyperparameters
    batch_size = 8
    seq_len    = 10

    # reuse from before
    transform = transforms.Compose([
        transforms.Resize((224,224)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.5]*3, std=[0.5]*3)
    ])

```

```

]) #dataset and transforms
train_ds = SonarSequenceDataset('train_split.csv', base_dir, seq_len, transform)
val_ds   = SonarSequenceDataset('val_split.csv',   base_dir, seq_len, transform)

# DataLoaders
train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True,
    ↪num_workers=4)
val_loader   = DataLoader(val_ds,   batch_size=batch_size, shuffle=False,
    ↪num_workers=4)

```

```

[ ]: import torch.optim as optim
import torch.nn as nn

# Instantiate model, loss, optimizer
model      = SonarCNNLSTM()
criterion = nn.MSELoss() # since speed, cog, depth are continuous
optimizer = optim.Adam(model.parameters(), lr=1e-3)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")
model.to(device)

# Training & Validation
num_epochs = 100
best_val_loss = float('inf')
train_losses, val_losses = [], []

for epoch in range(1, num_epochs+1):
    # - Training -
    model.train()
    running_train_loss = 0.0
    for batch in train_loader:
        imgs = batch['images'].to(device)    # (B, T, C, H, W)
        targets = torch.stack([batch['speed'],
                                batch['cog'],
                                batch['depth']], dim=1).to(device) # (B,3)

        preds = model(imgs)                  # (B,3)
        loss = criterion(preds, targets)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        running_train_loss += loss.item() * imgs.size(0)

    epoch_train_loss = running_train_loss / len(train_ds)

```



```

train_losses.append(epoch_train_loss)

# - Validation -
model.eval()
running_val_loss = 0.0
with torch.no_grad():
    for batch in val_loader:
        imgs      = batch['images'].to(device)
        targets   = torch.stack([batch['speed'],
                                batch['cog'],
                                batch['depth']], dim=1).to(device)

        preds     = model(imgs)
        loss       = criterion(preds, targets)
        running_val_loss += loss.item() * imgs.size(0)

epoch_val_loss = running_val_loss / len(val_ds)
val_losses.append(epoch_val_loss)

print(f"Epoch {epoch}/{num_epochs} - "
      f"Train Loss: {epoch_train_loss:.4f}  "
      f"Val Loss: {epoch_val_loss:.4f}")

# Save best model
if epoch_val_loss < best_val_loss:
    best_val_loss = epoch_val_loss
    torch.save(model.state_dict(), 'best_model001newfile.pth')

```

Using device: cuda

```

Epoch 1/100 - Train Loss: 1.9004  Val Loss: 0.1175
Epoch 2/100 - Train Loss: 0.9845  Val Loss: 0.2491
Epoch 3/100 - Train Loss: 0.6328  Val Loss: 0.2520
Epoch 4/100 - Train Loss: 0.5000  Val Loss: 0.2760
Epoch 5/100 - Train Loss: 0.3951  Val Loss: 0.2401
Epoch 6/100 - Train Loss: 0.4086  Val Loss: 0.6317
Epoch 7/100 - Train Loss: 0.4040  Val Loss: 0.2050
Epoch 8/100 - Train Loss: 0.3715  Val Loss: 0.2693
Epoch 9/100 - Train Loss: 0.3495  Val Loss: 0.3424
Epoch 10/100 - Train Loss: 0.3300  Val Loss: 0.3190
Epoch 11/100 - Train Loss: 0.3180  Val Loss: 0.1648
Epoch 12/100 - Train Loss: 0.3451  Val Loss: 0.3128
Epoch 13/100 - Train Loss: 0.3283  Val Loss: 0.1288
Epoch 14/100 - Train Loss: 0.3095  Val Loss: 0.0781
Epoch 15/100 - Train Loss: 0.3134  Val Loss: 0.0667
Epoch 16/100 - Train Loss: 0.3214  Val Loss: 0.0867
Epoch 17/100 - Train Loss: 0.3086  Val Loss: 0.0714
Epoch 18/100 - Train Loss: 0.3107  Val Loss: 0.1244
Epoch 19/100 - Train Loss: 0.3093  Val Loss: 0.0750

```

Epoch 20/100	- Train Loss: 0.2982	Val Loss: 0.0671
Epoch 21/100	- Train Loss: 0.3107	Val Loss: 0.1379
Epoch 22/100	- Train Loss: 0.3098	Val Loss: 0.1215
Epoch 23/100	- Train Loss: 0.3007	Val Loss: 0.0882
Epoch 24/100	- Train Loss: 0.3079	Val Loss: 0.1541
Epoch 25/100	- Train Loss: 0.2986	Val Loss: 0.1438
Epoch 26/100	- Train Loss: 0.3554	Val Loss: 0.1711
Epoch 27/100	- Train Loss: 0.3147	Val Loss: 0.0895
Epoch 28/100	- Train Loss: 0.2998	Val Loss: 0.2414
Epoch 29/100	- Train Loss: 0.3005	Val Loss: 0.1191
Epoch 30/100	- Train Loss: 0.2978	Val Loss: 0.0761
Epoch 31/100	- Train Loss: 0.2980	Val Loss: 0.0869
Epoch 32/100	- Train Loss: 0.3222	Val Loss: 0.1129
Epoch 33/100	- Train Loss: 0.2970	Val Loss: 0.0702
Epoch 34/100	- Train Loss: 0.2938	Val Loss: 0.0768
Epoch 35/100	- Train Loss: 0.3000	Val Loss: 0.0751
Epoch 36/100	- Train Loss: 0.2958	Val Loss: 0.1115
Epoch 37/100	- Train Loss: 0.2890	Val Loss: 0.1437
Epoch 38/100	- Train Loss: 0.2905	Val Loss: 0.0849
Epoch 39/100	- Train Loss: 0.3003	Val Loss: 0.0704
Epoch 40/100	- Train Loss: 0.2867	Val Loss: 0.0773
Epoch 41/100	- Train Loss: 0.2778	Val Loss: 0.0675
Epoch 42/100	- Train Loss: 0.2944	Val Loss: 0.0736
Epoch 43/100	- Train Loss: 0.2882	Val Loss: 0.0720
Epoch 44/100	- Train Loss: 0.2890	Val Loss: 0.3353
Epoch 45/100	- Train Loss: 0.3130	Val Loss: 0.0681
Epoch 46/100	- Train Loss: 0.2928	Val Loss: 0.0822
Epoch 47/100	- Train Loss: 0.2796	Val Loss: 0.0867
Epoch 48/100	- Train Loss: 0.2860	Val Loss: 0.0789
Epoch 49/100	- Train Loss: 0.2809	Val Loss: 0.0798
Epoch 50/100	- Train Loss: 0.2836	Val Loss: 0.0661
Epoch 51/100	- Train Loss: 0.2811	Val Loss: 0.0871
Epoch 52/100	- Train Loss: 0.2826	Val Loss: 0.0899
Epoch 53/100	- Train Loss: 0.2901	Val Loss: 0.0627
Epoch 54/100	- Train Loss: 0.2828	Val Loss: 0.0755
Epoch 55/100	- Train Loss: 0.2775	Val Loss: 0.0660
Epoch 56/100	- Train Loss: 0.2757	Val Loss: 0.1138
Epoch 57/100	- Train Loss: 0.2981	Val Loss: 0.0885
Epoch 58/100	- Train Loss: 0.2748	Val Loss: 0.0902
Epoch 59/100	- Train Loss: 0.2888	Val Loss: 0.1857
Epoch 60/100	- Train Loss: 0.2882	Val Loss: 0.0786
Epoch 61/100	- Train Loss: 0.2843	Val Loss: 0.0761
Epoch 62/100	- Train Loss: 0.2869	Val Loss: 0.0602
Epoch 63/100	- Train Loss: 0.2823	Val Loss: 0.1822
Epoch 64/100	- Train Loss: 0.2802	Val Loss: 0.1651
Epoch 65/100	- Train Loss: 0.2896	Val Loss: 0.0857
Epoch 66/100	- Train Loss: 0.2817	Val Loss: 0.0943
Epoch 67/100	- Train Loss: 0.2780	Val Loss: 0.1074

```

Epoch 68/100 - Train Loss: 0.2857 Val Loss: 0.0847
Epoch 69/100 - Train Loss: 0.2740 Val Loss: 0.0699
Epoch 70/100 - Train Loss: 0.2733 Val Loss: 0.1448
Epoch 71/100 - Train Loss: 0.2760 Val Loss: 0.0963
Epoch 72/100 - Train Loss: 0.2742 Val Loss: 0.0794
Epoch 73/100 - Train Loss: 0.2716 Val Loss: 0.0839
Epoch 74/100 - Train Loss: 0.2725 Val Loss: 0.0925
Epoch 75/100 - Train Loss: 0.2680 Val Loss: 0.0890
Epoch 76/100 - Train Loss: 0.2673 Val Loss: 0.1958
Epoch 77/100 - Train Loss: 0.2770 Val Loss: 0.0638
Epoch 78/100 - Train Loss: 0.2822 Val Loss: 0.1173
Epoch 79/100 - Train Loss: 0.2809 Val Loss: 0.0919
Epoch 80/100 - Train Loss: 0.2816 Val Loss: 0.1032
Epoch 81/100 - Train Loss: 0.2782 Val Loss: 0.0738
Epoch 82/100 - Train Loss: 0.2714 Val Loss: 0.0664
Epoch 83/100 - Train Loss: 0.2695 Val Loss: 0.1390
Epoch 84/100 - Train Loss: 0.2667 Val Loss: 0.1116
Epoch 85/100 - Train Loss: 0.2745 Val Loss: 0.0758
Epoch 86/100 - Train Loss: 0.2715 Val Loss: 0.0868
Epoch 87/100 - Train Loss: 0.2708 Val Loss: 0.0828
Epoch 88/100 - Train Loss: 0.2737 Val Loss: 0.0682
Epoch 89/100 - Train Loss: 0.2708 Val Loss: 0.0735
Epoch 90/100 - Train Loss: 0.2770 Val Loss: 0.0671
Epoch 91/100 - Train Loss: 0.2740 Val Loss: 0.1201
Epoch 92/100 - Train Loss: 0.2673 Val Loss: 0.0792
Epoch 93/100 - Train Loss: 0.2697 Val Loss: 0.0673
Epoch 94/100 - Train Loss: 0.2724 Val Loss: 0.0608
Epoch 95/100 - Train Loss: 0.2707 Val Loss: 0.0813
Epoch 96/100 - Train Loss: 0.2750 Val Loss: 0.0825
Epoch 97/100 - Train Loss: 0.2615 Val Loss: 0.1074
Epoch 98/100 - Train Loss: 0.2692 Val Loss: 0.0856
Epoch 99/100 - Train Loss: 0.2697 Val Loss: 0.1317
Epoch 100/100 - Train Loss: 0.2911 Val Loss: 0.0754


```

```

[10]: import matplotlib.pyplot as plt

epochs = range(1, num_epochs+1)
plt.plot(epochs, train_losses, label='Train Loss')
plt.plot(epochs, val_losses, label='Val Loss')
plt.xlabel('Epoch')
plt.ylabel('MSE Loss')
plt.legend()
plt.title('Training & Validation Loss')
plt.show()

```



```
[ ]: import torch
import torch.nn as nn
from torch.utils.data import DataLoader
from torchvision import transforms
from PIL import Image
import os
import pandas as pd

class SonarSequenceDataset(torch.utils.data.Dataset):
    def __init__(self, csv_path, base_dir, seq_len=10, transform=None):
        self.df = pd.read_csv(csv_path)
        self.base_dir = base_dir
        self.seq_len = seq_len
        self.transform = transform
        self.starts = list(range(len(self.df) - seq_len + 1))

    def __len__(self):
        return len(self.starts)

    def __getitem__(self, idx):
        start = self.starts[idx]
        records = self.df.iloc[start:start+self.seq_len]
        imgs = []
```

```

    for _, row in records.iterrows():
        path = os.path.join(self.base_dir, row['subset'], row['filepath'])
        img = Image.open(path).convert('RGB')
        if self.transform:
            img = self.transform(img)
        imgs.append(img)
    imgs = torch.stack(imgs, dim=0)
    last = records.iloc[-1]
    targets = torch.tensor([last['sog_speed'], last['cog_speed'],
↪last['waterdepth']], dtype=torch.float32)
    return {'images': imgs, 'targets': targets}

class SonarCNNLSTM(nn.Module):
    def __init__(self, cnn_out_dim=128, lstm_hidden=64, lstm_layers=1,
↪dropout=0.2, bidirectional=False, num_targets=3):
        super().__init__()
        from torchvision.models import resnet18
        backbone = resnet18(pretrained=False)
        in_features = backbone.fc.in_features
        backbone.fc = nn.Linear(in_features, cnn_out_dim)
        self.cnn = backbone
        self.lstm = nn.LSTM(
            input_size=cnn_out_dim,
            hidden_size=lstm_hidden,
            num_layers=lstm_layers,
            batch_first=True,
            dropout=dropout if lstm_layers > 1 else 0.0,
            bidirectional=bidirectional
        )
        lstm_dim = lstm_hidden * (2 if bidirectional else 1)
        self.regressor = nn.Sequential(
            nn.Dropout(dropout),
            nn.Linear(lstm_dim, lstm_dim // 2),
            nn.ReLU(inplace=True),
            nn.Linear(lstm_dim // 2, num_targets)
        )

    def forward(self, x):
        B, T, C, H, W = x.shape
        x = x.view(B * T, C, H, W)
        feats = self.cnn(x).view(B, T, -1)
        lstm_out, _ = self.lstm(feats)
        return self.regressor(lstm_out[:, -1, :])

# Load Test Set
transform = transforms.Compose([
    transforms.Resize((224, 224)),

```

```

        transforms.ToTensor(),
        transforms.Normalize(mean=[0.5]*3, std=[0.5]*3)
    ])
    base_dir = '/home/user/persistent/advanceaimethodsallfilestobeextracted22222222'
    test_ds = SonarSequenceDataset('test_split.csv', base_dir, seq_len=10,
        ↪transform=transform)
    test_loader = DataLoader(test_ds, batch_size=8, shuffle=False)

    # Evaluate
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    model = SonarCNNLSTM()
    model.load_state_dict(torch.load('best_model001newfile.pth',
        ↪map_location=device))
    model.to(device).eval()

    criterion = nn.MSELoss(reduction='none')
    all_losses, all_targets, all_preds = [], [], []

    with torch.no_grad():
        for batch in test_loader:
            imgs = batch['images'].to(device)
            targets = batch['targets'].to(device)
            preds = model(imgs)
            loss = criterion(preds, targets)
            all_losses.append(loss.cpu())
            all_targets.append(targets.cpu())
            all_preds.append(preds.cpu())

    losses = torch.cat(all_losses, dim=0).numpy()
    targets = torch.cat(all_targets, dim=0).numpy()
    preds = torch.cat(all_preds, dim=0).numpy()

    mse_per_target = losses.mean(axis=0)
    overall_mse = mse_per_target.mean()

    df_metrics = pd.DataFrame({
        'Target': ['Speed', 'COG', 'Depth', 'Overall'],
        'MSE': [mse_per_target[0], mse_per_target[1], mse_per_target[2],
        ↪overall_mse]
    })

    print("\nTest Evaluation Results (Mean Squared Error per target):")
    print(df_metrics)

```

Test Evaluation Results (Mean Squared Error per target):

Target	MSE
--------	-----

```
0    Speed    0.040594
1     COG    12.258199
2    Depth    4.138172
3 Overall    5.478988
```

```
[12]: def train_model(run_id, model, train_loader, val_loader, criterion, optimizer,
    ↪ num_epochs=100, device='cpu'):
    train_losses, val_losses = [], []
    best_model_wts = None
    best_val_loss = float('inf')

    for epoch in range(num_epochs):
        # Training
        model.train()
        train_loss = 0.0
        for batch in train_loader:
            imgs = batch['images'].to(device)
            targets = torch.stack([batch['speed'], batch['cog'],
    ↪ batch['depth']], dim=1).to(device)

            preds = model(imgs)
            loss = criterion(preds, targets)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()
            train_loss += loss.item() * imgs.size(0)

        epoch_train_loss = train_loss / len(train_loader.dataset)
        train_losses.append(epoch_train_loss)

        # Validation
        model.eval()
        val_loss = 0.0
        with torch.no_grad():
            for batch in val_loader:
                imgs = batch['images'].to(device)
                targets = torch.stack([batch['speed'], batch['cog'],
    ↪ batch['depth']], dim=1).to(device)

                preds = model(imgs)
                loss = criterion(preds, targets)
                val_loss += loss.item() * imgs.size(0)

        epoch_val_loss = val_loss / len(val_loader.dataset)
        val_losses.append(epoch_val_loss)
```

```

        print(f"Run {run_id} | Epoch {epoch+1}/{num_epochs} - Train Loss:␣
→{epoch_train_loss:.4f} | Val Loss: {epoch_val_loss:.4f}")

        # Save best
        if epoch_val_loss < best_val_loss:
            best_val_loss = epoch_val_loss
            best_model_wts = model.state_dict()

        # Save best weights
        torch.save(best_model_wts, f'model_run{run_id}.pt')
    return train_losses, val_losses

```

```

[13]: import torch.nn as nn
import torch.optim as optim

all_losses = []

for run in range(1, 4):
    print(f"\n=== Training Run {run}/3 ===\n")
    model = SonarCNNLSTM().to(device)
    criterion = nn.MSELoss()
    optimizer = optim.Adam(model.parameters(), lr=1e-4)

    train_losses, val_losses = train_model(
        run_id=run,
        model=model,
        train_loader=train_loader,
        val_loader=val_loader,
        criterion=criterion,
        optimizer=optimizer,
        num_epochs=100,
        device=device
    )
    all_losses.append((train_losses, val_losses))

```

=== Training Run 1/3 ===

```

Run 1 | Epoch 1/100 - Train Loss: 4.5290 | Val Loss: 0.7208
Run 1 | Epoch 2/100 - Train Loss: 1.1570 | Val Loss: 0.2657
Run 1 | Epoch 3/100 - Train Loss: 0.5527 | Val Loss: 0.1624
Run 1 | Epoch 4/100 - Train Loss: 0.4156 | Val Loss: 0.1445
Run 1 | Epoch 5/100 - Train Loss: 0.3340 | Val Loss: 0.1006
Run 1 | Epoch 6/100 - Train Loss: 0.2752 | Val Loss: 0.0864
Run 1 | Epoch 7/100 - Train Loss: 0.2622 | Val Loss: 0.0751
Run 1 | Epoch 8/100 - Train Loss: 0.2728 | Val Loss: 0.1092
Run 1 | Epoch 9/100 - Train Loss: 0.2343 | Val Loss: 0.0731

```


Run 1		Epoch 10/100	-	Train Loss: 0.2028		Val Loss: 0.0816
Run 1		Epoch 11/100	-	Train Loss: 0.2092		Val Loss: 0.0692
Run 1		Epoch 12/100	-	Train Loss: 0.2015		Val Loss: 0.0739
Run 1		Epoch 13/100	-	Train Loss: 0.2004		Val Loss: 0.0805
Run 1		Epoch 14/100	-	Train Loss: 0.1975		Val Loss: 0.0676
Run 1		Epoch 15/100	-	Train Loss: 0.1565		Val Loss: 0.0692
Run 1		Epoch 16/100	-	Train Loss: 0.1331		Val Loss: 0.0692
Run 1		Epoch 17/100	-	Train Loss: 0.1202		Val Loss: 0.0693
Run 1		Epoch 18/100	-	Train Loss: 0.1174		Val Loss: 0.0751
Run 1		Epoch 19/100	-	Train Loss: 0.1123		Val Loss: 0.0864
Run 1		Epoch 20/100	-	Train Loss: 0.1164		Val Loss: 0.0720
Run 1		Epoch 21/100	-	Train Loss: 0.1092		Val Loss: 0.0708
Run 1		Epoch 22/100	-	Train Loss: 0.1142		Val Loss: 0.0791
Run 1		Epoch 23/100	-	Train Loss: 0.1091		Val Loss: 0.0751
Run 1		Epoch 24/100	-	Train Loss: 0.1110		Val Loss: 0.0712
Run 1		Epoch 25/100	-	Train Loss: 0.1092		Val Loss: 0.0784
Run 1		Epoch 26/100	-	Train Loss: 0.1014		Val Loss: 0.0759
Run 1		Epoch 27/100	-	Train Loss: 0.1263		Val Loss: 0.0709
Run 1		Epoch 28/100	-	Train Loss: 0.0972		Val Loss: 0.0713
Run 1		Epoch 29/100	-	Train Loss: 0.0984		Val Loss: 0.0797
Run 1		Epoch 30/100	-	Train Loss: 0.1117		Val Loss: 0.0720
Run 1		Epoch 31/100	-	Train Loss: 0.0909		Val Loss: 0.0928
Run 1		Epoch 32/100	-	Train Loss: 0.0843		Val Loss: 0.0934
Run 1		Epoch 33/100	-	Train Loss: 0.0868		Val Loss: 0.1001
Run 1		Epoch 34/100	-	Train Loss: 0.0805		Val Loss: 0.1073
Run 1		Epoch 35/100	-	Train Loss: 0.0947		Val Loss: 0.0805
Run 1		Epoch 36/100	-	Train Loss: 0.0925		Val Loss: 0.0707
Run 1		Epoch 37/100	-	Train Loss: 0.0843		Val Loss: 0.0810
Run 1		Epoch 38/100	-	Train Loss: 0.0889		Val Loss: 0.0771
Run 1		Epoch 39/100	-	Train Loss: 0.0839		Val Loss: 0.0833
Run 1		Epoch 40/100	-	Train Loss: 0.0890		Val Loss: 0.1025
Run 1		Epoch 41/100	-	Train Loss: 0.0844		Val Loss: 0.0745
Run 1		Epoch 42/100	-	Train Loss: 0.0798		Val Loss: 0.0856
Run 1		Epoch 43/100	-	Train Loss: 0.0764		Val Loss: 0.0738
Run 1		Epoch 44/100	-	Train Loss: 0.0735		Val Loss: 0.0733
Run 1		Epoch 45/100	-	Train Loss: 0.0732		Val Loss: 0.0851
Run 1		Epoch 46/100	-	Train Loss: 0.0724		Val Loss: 0.0815
Run 1		Epoch 47/100	-	Train Loss: 0.0847		Val Loss: 0.0676
Run 1		Epoch 48/100	-	Train Loss: 0.0724		Val Loss: 0.0728
Run 1		Epoch 49/100	-	Train Loss: 0.0664		Val Loss: 0.0925
Run 1		Epoch 50/100	-	Train Loss: 0.0769		Val Loss: 0.0720
Run 1		Epoch 51/100	-	Train Loss: 0.0720		Val Loss: 0.0997
Run 1		Epoch 52/100	-	Train Loss: 0.0721		Val Loss: 0.0760
Run 1		Epoch 53/100	-	Train Loss: 0.0658		Val Loss: 0.0813
Run 1		Epoch 54/100	-	Train Loss: 0.0721		Val Loss: 0.0942
Run 1		Epoch 55/100	-	Train Loss: 0.0750		Val Loss: 0.0852
Run 1		Epoch 56/100	-	Train Loss: 0.0671		Val Loss: 0.0750
Run 1		Epoch 57/100	-	Train Loss: 0.0695		Val Loss: 0.0698

Run 1 | Epoch 58/100 - Train Loss: 0.0666 | Val Loss: 0.0747
Run 1 | Epoch 59/100 - Train Loss: 0.0622 | Val Loss: 0.0738
Run 1 | Epoch 60/100 - Train Loss: 0.0735 | Val Loss: 0.0794
Run 1 | Epoch 61/100 - Train Loss: 0.0820 | Val Loss: 0.0681
Run 1 | Epoch 62/100 - Train Loss: 0.0696 | Val Loss: 0.0943
Run 1 | Epoch 63/100 - Train Loss: 0.0620 | Val Loss: 0.0911
Run 1 | Epoch 64/100 - Train Loss: 0.0592 | Val Loss: 0.0736
Run 1 | Epoch 65/100 - Train Loss: 0.0610 | Val Loss: 0.0734
Run 1 | Epoch 66/100 - Train Loss: 0.0629 | Val Loss: 0.0832
Run 1 | Epoch 67/100 - Train Loss: 0.0719 | Val Loss: 0.0664
Run 1 | Epoch 68/100 - Train Loss: 0.0666 | Val Loss: 0.0762
Run 1 | Epoch 69/100 - Train Loss: 0.0611 | Val Loss: 0.0766
Run 1 | Epoch 70/100 - Train Loss: 0.0593 | Val Loss: 0.0719
Run 1 | Epoch 71/100 - Train Loss: 0.0579 | Val Loss: 0.0801
Run 1 | Epoch 72/100 - Train Loss: 0.0618 | Val Loss: 0.0769
Run 1 | Epoch 73/100 - Train Loss: 0.0584 | Val Loss: 0.0751
Run 1 | Epoch 74/100 - Train Loss: 0.0549 | Val Loss: 0.0845
Run 1 | Epoch 75/100 - Train Loss: 0.0565 | Val Loss: 0.0842
Run 1 | Epoch 76/100 - Train Loss: 0.0567 | Val Loss: 0.0884
Run 1 | Epoch 77/100 - Train Loss: 0.0555 | Val Loss: 0.0836
Run 1 | Epoch 78/100 - Train Loss: 0.0539 | Val Loss: 0.0698
Run 1 | Epoch 79/100 - Train Loss: 0.0600 | Val Loss: 0.0698
Run 1 | Epoch 80/100 - Train Loss: 0.0523 | Val Loss: 0.0748
Run 1 | Epoch 81/100 - Train Loss: 0.0582 | Val Loss: 0.1009
Run 1 | Epoch 82/100 - Train Loss: 0.0529 | Val Loss: 0.0691
Run 1 | Epoch 83/100 - Train Loss: 0.0522 | Val Loss: 0.0728
Run 1 | Epoch 84/100 - Train Loss: 0.0550 | Val Loss: 0.0731
Run 1 | Epoch 85/100 - Train Loss: 0.0523 | Val Loss: 0.0840
Run 1 | Epoch 86/100 - Train Loss: 0.0538 | Val Loss: 0.0910
Run 1 | Epoch 87/100 - Train Loss: 0.0516 | Val Loss: 0.0827
Run 1 | Epoch 88/100 - Train Loss: 0.0921 | Val Loss: 0.0874
Run 1 | Epoch 89/100 - Train Loss: 0.0722 | Val Loss: 0.0857
Run 1 | Epoch 90/100 - Train Loss: 0.0570 | Val Loss: 0.0660
Run 1 | Epoch 91/100 - Train Loss: 0.0665 | Val Loss: 0.0643
Run 1 | Epoch 92/100 - Train Loss: 0.0501 | Val Loss: 0.0727
Run 1 | Epoch 93/100 - Train Loss: 0.0522 | Val Loss: 0.0742
Run 1 | Epoch 94/100 - Train Loss: 0.0543 | Val Loss: 0.0705
Run 1 | Epoch 95/100 - Train Loss: 0.0504 | Val Loss: 0.0864
Run 1 | Epoch 96/100 - Train Loss: 0.0479 | Val Loss: 0.0675
Run 1 | Epoch 97/100 - Train Loss: 0.0452 | Val Loss: 0.0845
Run 1 | Epoch 98/100 - Train Loss: 0.0479 | Val Loss: 0.0754
Run 1 | Epoch 99/100 - Train Loss: 0.0525 | Val Loss: 0.0727
Run 1 | Epoch 100/100 - Train Loss: 0.0447 | Val Loss: 0.0900

=== Training Run 2/3 ===

Run 2 | Epoch 1/100 - Train Loss: 5.6752 | Val Loss: 1.0592
Run 2 | Epoch 2/100 - Train Loss: 1.6040 | Val Loss: 0.0655

Run 2		Epoch 3/100	-	Train Loss: 0.6318		Val Loss: 0.1318
Run 2		Epoch 4/100	-	Train Loss: 0.3678		Val Loss: 0.1818
Run 2		Epoch 5/100	-	Train Loss: 0.2955		Val Loss: 0.1105
Run 2		Epoch 6/100	-	Train Loss: 0.3134		Val Loss: 0.1310
Run 2		Epoch 7/100	-	Train Loss: 0.2570		Val Loss: 0.1079
Run 2		Epoch 8/100	-	Train Loss: 0.2203		Val Loss: 0.0575
Run 2		Epoch 9/100	-	Train Loss: 0.2052		Val Loss: 0.0656
Run 2		Epoch 10/100	-	Train Loss: 0.1575		Val Loss: 0.0637
Run 2		Epoch 11/100	-	Train Loss: 0.1686		Val Loss: 0.0743
Run 2		Epoch 12/100	-	Train Loss: 0.1860		Val Loss: 0.0741
Run 2		Epoch 13/100	-	Train Loss: 0.1257		Val Loss: 0.0778
Run 2		Epoch 14/100	-	Train Loss: 0.1226		Val Loss: 0.0814
Run 2		Epoch 15/100	-	Train Loss: 0.1239		Val Loss: 0.0751
Run 2		Epoch 16/100	-	Train Loss: 0.1293		Val Loss: 0.0956
Run 2		Epoch 17/100	-	Train Loss: 0.1399		Val Loss: 0.2234
Run 2		Epoch 18/100	-	Train Loss: 0.1437		Val Loss: 0.0727
Run 2		Epoch 19/100	-	Train Loss: 0.1084		Val Loss: 0.0733
Run 2		Epoch 20/100	-	Train Loss: 0.1084		Val Loss: 0.0768
Run 2		Epoch 21/100	-	Train Loss: 0.1085		Val Loss: 0.0748
Run 2		Epoch 22/100	-	Train Loss: 0.1040		Val Loss: 0.0730
Run 2		Epoch 23/100	-	Train Loss: 0.1149		Val Loss: 0.0863
Run 2		Epoch 24/100	-	Train Loss: 0.0981		Val Loss: 0.0771
Run 2		Epoch 25/100	-	Train Loss: 0.0888		Val Loss: 0.0668
Run 2		Epoch 26/100	-	Train Loss: 0.0899		Val Loss: 0.0725
Run 2		Epoch 27/100	-	Train Loss: 0.0876		Val Loss: 0.0725
Run 2		Epoch 28/100	-	Train Loss: 0.0927		Val Loss: 0.0681
Run 2		Epoch 29/100	-	Train Loss: 0.0928		Val Loss: 0.0812
Run 2		Epoch 30/100	-	Train Loss: 0.0996		Val Loss: 0.0788
Run 2		Epoch 31/100	-	Train Loss: 0.0916		Val Loss: 0.0754
Run 2		Epoch 32/100	-	Train Loss: 0.0824		Val Loss: 0.0806
Run 2		Epoch 33/100	-	Train Loss: 0.0742		Val Loss: 0.0804
Run 2		Epoch 34/100	-	Train Loss: 0.0837		Val Loss: 0.0772
Run 2		Epoch 35/100	-	Train Loss: 0.0816		Val Loss: 0.0760
Run 2		Epoch 36/100	-	Train Loss: 0.0805		Val Loss: 0.0786
Run 2		Epoch 37/100	-	Train Loss: 0.0883		Val Loss: 0.0766
Run 2		Epoch 38/100	-	Train Loss: 0.0772		Val Loss: 0.0772
Run 2		Epoch 39/100	-	Train Loss: 0.0808		Val Loss: 0.0768
Run 2		Epoch 40/100	-	Train Loss: 0.0764		Val Loss: 0.0834
Run 2		Epoch 41/100	-	Train Loss: 0.0725		Val Loss: 0.0833
Run 2		Epoch 42/100	-	Train Loss: 0.1070		Val Loss: 0.0822
Run 2		Epoch 43/100	-	Train Loss: 0.0910		Val Loss: 0.0846
Run 2		Epoch 44/100	-	Train Loss: 0.0833		Val Loss: 0.0879
Run 2		Epoch 45/100	-	Train Loss: 0.0773		Val Loss: 0.0730
Run 2		Epoch 46/100	-	Train Loss: 0.0775		Val Loss: 0.0800
Run 2		Epoch 47/100	-	Train Loss: 0.0808		Val Loss: 0.0813
Run 2		Epoch 48/100	-	Train Loss: 0.0680		Val Loss: 0.0915
Run 2		Epoch 49/100	-	Train Loss: 0.0742		Val Loss: 0.0741
Run 2		Epoch 50/100	-	Train Loss: 0.0704		Val Loss: 0.0973

Run 2	Epoch 51/100	- Train Loss: 0.0740	Val Loss: 0.0747
Run 2	Epoch 52/100	- Train Loss: 0.0665	Val Loss: 0.0769
Run 2	Epoch 53/100	- Train Loss: 0.0666	Val Loss: 0.0748
Run 2	Epoch 54/100	- Train Loss: 0.0668	Val Loss: 0.0745
Run 2	Epoch 55/100	- Train Loss: 0.0759	Val Loss: 0.0806
Run 2	Epoch 56/100	- Train Loss: 0.0696	Val Loss: 0.0964
Run 2	Epoch 57/100	- Train Loss: 0.0679	Val Loss: 0.0771
Run 2	Epoch 58/100	- Train Loss: 0.0716	Val Loss: 0.1074
Run 2	Epoch 59/100	- Train Loss: 0.0907	Val Loss: 0.0910
Run 2	Epoch 60/100	- Train Loss: 0.0719	Val Loss: 0.0707
Run 2	Epoch 61/100	- Train Loss: 0.0639	Val Loss: 0.0730
Run 2	Epoch 62/100	- Train Loss: 0.0694	Val Loss: 0.0730
Run 2	Epoch 63/100	- Train Loss: 0.0816	Val Loss: 0.0918
Run 2	Epoch 64/100	- Train Loss: 0.0694	Val Loss: 0.0864
Run 2	Epoch 65/100	- Train Loss: 0.0637	Val Loss: 0.0810
Run 2	Epoch 66/100	- Train Loss: 0.0612	Val Loss: 0.0695
Run 2	Epoch 67/100	- Train Loss: 0.0621	Val Loss: 0.0772
Run 2	Epoch 68/100	- Train Loss: 0.0604	Val Loss: 0.0762
Run 2	Epoch 69/100	- Train Loss: 0.0591	Val Loss: 0.0875
Run 2	Epoch 70/100	- Train Loss: 0.0634	Val Loss: 0.0712
Run 2	Epoch 71/100	- Train Loss: 0.0618	Val Loss: 0.0742
Run 2	Epoch 72/100	- Train Loss: 0.0591	Val Loss: 0.0727
Run 2	Epoch 73/100	- Train Loss: 0.0581	Val Loss: 0.0733
Run 2	Epoch 74/100	- Train Loss: 0.0553	Val Loss: 0.0708
Run 2	Epoch 75/100	- Train Loss: 0.0563	Val Loss: 0.0877
Run 2	Epoch 76/100	- Train Loss: 0.0593	Val Loss: 0.0729
Run 2	Epoch 77/100	- Train Loss: 0.0557	Val Loss: 0.0709
Run 2	Epoch 78/100	- Train Loss: 0.0703	Val Loss: 0.0686
Run 2	Epoch 79/100	- Train Loss: 0.0712	Val Loss: 0.0808
Run 2	Epoch 80/100	- Train Loss: 0.0592	Val Loss: 0.0722
Run 2	Epoch 81/100	- Train Loss: 0.0576	Val Loss: 0.0743
Run 2	Epoch 82/100	- Train Loss: 0.0558	Val Loss: 0.0722
Run 2	Epoch 83/100	- Train Loss: 0.0524	Val Loss: 0.0739
Run 2	Epoch 84/100	- Train Loss: 0.0526	Val Loss: 0.0742
Run 2	Epoch 85/100	- Train Loss: 0.0541	Val Loss: 0.0713
Run 2	Epoch 86/100	- Train Loss: 0.0547	Val Loss: 0.0718
Run 2	Epoch 87/100	- Train Loss: 0.0496	Val Loss: 0.0693
Run 2	Epoch 88/100	- Train Loss: 0.0531	Val Loss: 0.0719
Run 2	Epoch 89/100	- Train Loss: 0.0671	Val Loss: 0.0809
Run 2	Epoch 90/100	- Train Loss: 0.0541	Val Loss: 0.0702
Run 2	Epoch 91/100	- Train Loss: 0.0507	Val Loss: 0.0672
Run 2	Epoch 92/100	- Train Loss: 0.0486	Val Loss: 0.0718
Run 2	Epoch 93/100	- Train Loss: 0.0486	Val Loss: 0.0692
Run 2	Epoch 94/100	- Train Loss: 0.0528	Val Loss: 0.0653
Run 2	Epoch 95/100	- Train Loss: 0.0546	Val Loss: 0.0701
Run 2	Epoch 96/100	- Train Loss: 0.0490	Val Loss: 0.0788
Run 2	Epoch 97/100	- Train Loss: 0.0472	Val Loss: 0.0707
Run 2	Epoch 98/100	- Train Loss: 0.0452	Val Loss: 0.0672

Run 2 | Epoch 99/100 - Train Loss: 0.0461 | Val Loss: 0.0717
Run 2 | Epoch 100/100 - Train Loss: 0.0471 | Val Loss: 0.0688

=== Training Run 3/3 ===

Run 3 | Epoch 1/100 - Train Loss: 3.7300 | Val Loss: 0.1372
Run 3 | Epoch 2/100 - Train Loss: 0.9523 | Val Loss: 0.0724
Run 3 | Epoch 3/100 - Train Loss: 0.5008 | Val Loss: 0.0772
Run 3 | Epoch 4/100 - Train Loss: 0.3912 | Val Loss: 0.1128
Run 3 | Epoch 5/100 - Train Loss: 0.3035 | Val Loss: 0.1292
Run 3 | Epoch 6/100 - Train Loss: 0.2657 | Val Loss: 0.1489
Run 3 | Epoch 7/100 - Train Loss: 0.2312 | Val Loss: 0.1432
Run 3 | Epoch 8/100 - Train Loss: 0.2310 | Val Loss: 0.0910
Run 3 | Epoch 9/100 - Train Loss: 0.2187 | Val Loss: 0.1068
Run 3 | Epoch 10/100 - Train Loss: 0.2217 | Val Loss: 0.0678
Run 3 | Epoch 11/100 - Train Loss: 0.2153 | Val Loss: 0.0958
Run 3 | Epoch 12/100 - Train Loss: 0.1896 | Val Loss: 0.0976
Run 3 | Epoch 13/100 - Train Loss: 0.1450 | Val Loss: 0.0884
Run 3 | Epoch 14/100 - Train Loss: 0.1364 | Val Loss: 0.0811
Run 3 | Epoch 15/100 - Train Loss: 0.1257 | Val Loss: 0.0733
Run 3 | Epoch 16/100 - Train Loss: 0.1146 | Val Loss: 0.1054
Run 3 | Epoch 17/100 - Train Loss: 0.1109 | Val Loss: 0.0934
Run 3 | Epoch 18/100 - Train Loss: 0.1107 | Val Loss: 0.0774
Run 3 | Epoch 19/100 - Train Loss: 0.1150 | Val Loss: 0.0856
Run 3 | Epoch 20/100 - Train Loss: 0.1294 | Val Loss: 0.0920
Run 3 | Epoch 21/100 - Train Loss: 0.1050 | Val Loss: 0.0801
Run 3 | Epoch 22/100 - Train Loss: 0.0918 | Val Loss: 0.0906
Run 3 | Epoch 23/100 - Train Loss: 0.0936 | Val Loss: 0.0725
Run 3 | Epoch 24/100 - Train Loss: 0.1182 | Val Loss: 0.0778
Run 3 | Epoch 25/100 - Train Loss: 0.1109 | Val Loss: 0.0897
Run 3 | Epoch 26/100 - Train Loss: 0.0984 | Val Loss: 0.0854
Run 3 | Epoch 27/100 - Train Loss: 0.0893 | Val Loss: 0.0944
Run 3 | Epoch 28/100 - Train Loss: 0.0886 | Val Loss: 0.0775
Run 3 | Epoch 29/100 - Train Loss: 0.0987 | Val Loss: 0.0757
Run 3 | Epoch 30/100 - Train Loss: 0.0864 | Val Loss: 0.0820
Run 3 | Epoch 31/100 - Train Loss: 0.0800 | Val Loss: 0.0923
Run 3 | Epoch 32/100 - Train Loss: 0.0855 | Val Loss: 0.0878
Run 3 | Epoch 33/100 - Train Loss: 0.1021 | Val Loss: 0.0746
Run 3 | Epoch 34/100 - Train Loss: 0.0868 | Val Loss: 0.0764
Run 3 | Epoch 35/100 - Train Loss: 0.0789 | Val Loss: 0.0679
Run 3 | Epoch 36/100 - Train Loss: 0.0784 | Val Loss: 0.0686
Run 3 | Epoch 37/100 - Train Loss: 0.0829 | Val Loss: 0.0743
Run 3 | Epoch 38/100 - Train Loss: 0.0967 | Val Loss: 0.0715
Run 3 | Epoch 39/100 - Train Loss: 0.1130 | Val Loss: 0.0685
Run 3 | Epoch 40/100 - Train Loss: 0.0866 | Val Loss: 0.0721
Run 3 | Epoch 41/100 - Train Loss: 0.0869 | Val Loss: 0.0687
Run 3 | Epoch 42/100 - Train Loss: 0.0823 | Val Loss: 0.0680
Run 3 | Epoch 43/100 - Train Loss: 0.0743 | Val Loss: 0.0820

Run 3		Epoch 44/100	-	Train Loss: 0.0748		Val Loss: 0.0718
Run 3		Epoch 45/100	-	Train Loss: 0.0773		Val Loss: 0.0755
Run 3		Epoch 46/100	-	Train Loss: 0.0691		Val Loss: 0.0702
Run 3		Epoch 47/100	-	Train Loss: 0.0692		Val Loss: 0.0747
Run 3		Epoch 48/100	-	Train Loss: 0.0721		Val Loss: 0.0714
Run 3		Epoch 49/100	-	Train Loss: 0.1200		Val Loss: 0.0839
Run 3		Epoch 50/100	-	Train Loss: 0.0740		Val Loss: 0.0770
Run 3		Epoch 51/100	-	Train Loss: 0.0653		Val Loss: 0.0704
Run 3		Epoch 52/100	-	Train Loss: 0.0761		Val Loss: 0.0810
Run 3		Epoch 53/100	-	Train Loss: 0.0797		Val Loss: 0.0741
Run 3		Epoch 54/100	-	Train Loss: 0.0827		Val Loss: 0.0721
Run 3		Epoch 55/100	-	Train Loss: 0.0651		Val Loss: 0.0751
Run 3		Epoch 56/100	-	Train Loss: 0.0680		Val Loss: 0.0862
Run 3		Epoch 57/100	-	Train Loss: 0.0633		Val Loss: 0.0711
Run 3		Epoch 58/100	-	Train Loss: 0.0624		Val Loss: 0.0746
Run 3		Epoch 59/100	-	Train Loss: 0.0688		Val Loss: 0.0693
Run 3		Epoch 60/100	-	Train Loss: 0.0671		Val Loss: 0.0773
Run 3		Epoch 61/100	-	Train Loss: 0.0813		Val Loss: 0.0826
Run 3		Epoch 62/100	-	Train Loss: 0.0728		Val Loss: 0.0775
Run 3		Epoch 63/100	-	Train Loss: 0.0660		Val Loss: 0.0651
Run 3		Epoch 64/100	-	Train Loss: 0.0615		Val Loss: 0.0725
Run 3		Epoch 65/100	-	Train Loss: 0.0631		Val Loss: 0.0686
Run 3		Epoch 66/100	-	Train Loss: 0.0567		Val Loss: 0.0987
Run 3		Epoch 67/100	-	Train Loss: 0.0564		Val Loss: 0.0813
Run 3		Epoch 68/100	-	Train Loss: 0.0583		Val Loss: 0.0721
Run 3		Epoch 69/100	-	Train Loss: 0.0542		Val Loss: 0.0785
Run 3		Epoch 70/100	-	Train Loss: 0.0577		Val Loss: 0.0801
Run 3		Epoch 71/100	-	Train Loss: 0.0534		Val Loss: 0.0678
Run 3		Epoch 72/100	-	Train Loss: 0.0538		Val Loss: 0.0802
Run 3		Epoch 73/100	-	Train Loss: 0.0632		Val Loss: 0.0764
Run 3		Epoch 74/100	-	Train Loss: 0.0655		Val Loss: 0.0959
Run 3		Epoch 75/100	-	Train Loss: 0.0603		Val Loss: 0.0648
Run 3		Epoch 76/100	-	Train Loss: 0.0588		Val Loss: 0.0659
Run 3		Epoch 77/100	-	Train Loss: 0.0577		Val Loss: 0.0990
Run 3		Epoch 78/100	-	Train Loss: 0.0542		Val Loss: 0.0706
Run 3		Epoch 79/100	-	Train Loss: 0.0525		Val Loss: 0.0929
Run 3		Epoch 80/100	-	Train Loss: 0.0536		Val Loss: 0.0723
Run 3		Epoch 81/100	-	Train Loss: 0.0507		Val Loss: 0.0677
Run 3		Epoch 82/100	-	Train Loss: 0.0505		Val Loss: 0.0912
Run 3		Epoch 83/100	-	Train Loss: 0.0494		Val Loss: 0.0733
Run 3		Epoch 84/100	-	Train Loss: 0.0529		Val Loss: 0.0770
Run 3		Epoch 85/100	-	Train Loss: 0.0497		Val Loss: 0.0921
Run 3		Epoch 86/100	-	Train Loss: 0.0528		Val Loss: 0.0712
Run 3		Epoch 87/100	-	Train Loss: 0.0565		Val Loss: 0.0691
Run 3		Epoch 88/100	-	Train Loss: 0.0633		Val Loss: 0.0709
Run 3		Epoch 89/100	-	Train Loss: 0.0497		Val Loss: 0.0835
Run 3		Epoch 90/100	-	Train Loss: 0.0498		Val Loss: 0.0800
Run 3		Epoch 91/100	-	Train Loss: 0.0495		Val Loss: 0.0797

```

Run 3 | Epoch 92/100 - Train Loss: 0.0493 | Val Loss: 0.0707
Run 3 | Epoch 93/100 - Train Loss: 0.0470 | Val Loss: 0.0710
Run 3 | Epoch 94/100 - Train Loss: 0.0459 | Val Loss: 0.0919
Run 3 | Epoch 95/100 - Train Loss: 0.0561 | Val Loss: 0.0686
Run 3 | Epoch 96/100 - Train Loss: 0.0480 | Val Loss: 0.0709
Run 3 | Epoch 97/100 - Train Loss: 0.0480 | Val Loss: 0.0643
Run 3 | Epoch 98/100 - Train Loss: 0.0475 | Val Loss: 0.0775
Run 3 | Epoch 99/100 - Train Loss: 0.0468 | Val Loss: 0.0745
Run 3 | Epoch 100/100 - Train Loss: 0.0451 | Val Loss: 0.0741

```

```

[16]: def evaluate_model(model_path, test_loader):
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    model = SonarCNNLSTM()
    model.load_state_dict(torch.load(model_path, map_location=device))
    model.to(device).eval()

    criterion = nn.MSELoss(reduction='none')
    all_losses = []

    with torch.no_grad():
        for batch in test_loader:
            imgs = batch['images'].to(device)
            targets = batch['targets'].to(device) # ← fix here
            outputs = model(imgs)
            loss = criterion(outputs, targets)
            all_losses.append(loss.cpu())

    losses = torch.cat(all_losses, dim=0).numpy()
    mse_per_target = losses.mean(axis=0)
    overall_mse = mse_per_target.mean()
    return mse_per_target, overall_mse

```

```

[17]: for i in range(1, 4):
    print(f"\nEvaluation of Run {i}")
    mse_vals, mse_overall = evaluate_model(f'model_run{i}.pt', test_loader)
    print(f"  Speed MSE: {mse_vals[0]:.4f}")
    print(f"  COG MSE:   {mse_vals[1]:.4f}")
    print(f"  Depth MSE: {mse_vals[2]:.4f}")
    print(f"  Overall MSE: {mse_overall:.4f}")

```

```

Evaluation of Run 1
  Speed MSE: 0.0969
  COG MSE:   7.3301
  Depth MSE: 6.3422
  Overall MSE: 4.5898

```

```

Evaluation of Run 2

```

Speed MSE: 0.0920
COG MSE: 6.5768
Depth MSE: 6.6812
Overall MSE: 4.4500

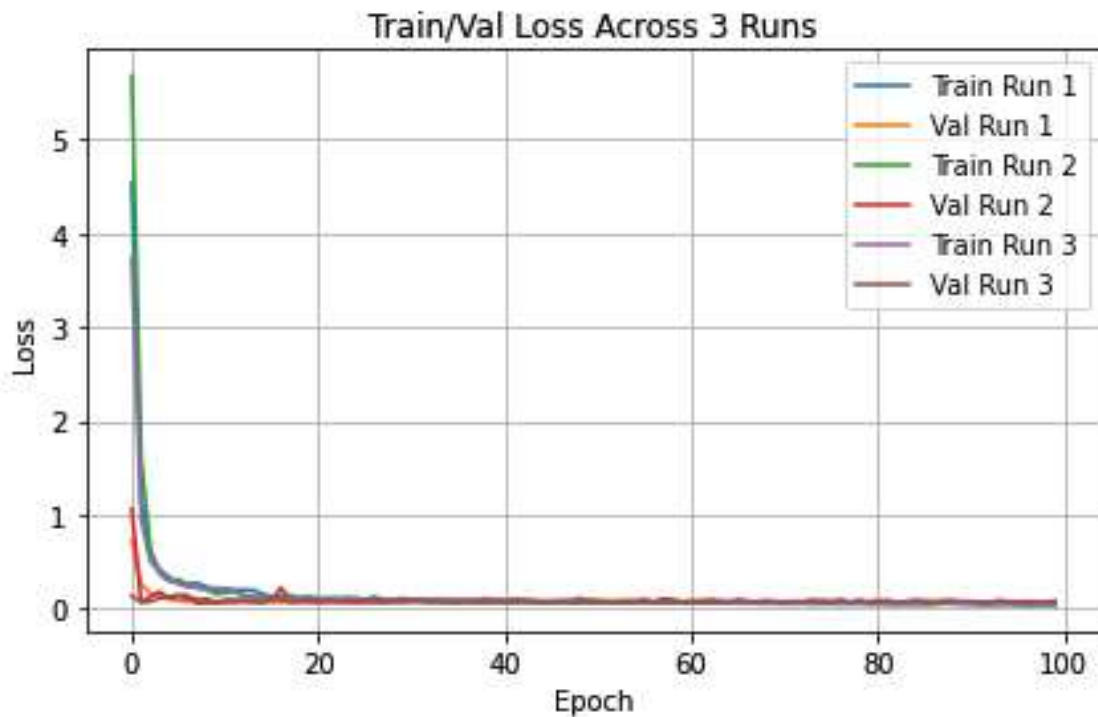
Evaluation of Run 3

Speed MSE: 0.1059
COG MSE: 6.4471
Depth MSE: 6.5358
Overall MSE: 4.3629

```
[18]: import matplotlib.pyplot as plt

for i, (train_loss, val_loss) in enumerate(all_losses, 1):
    plt.plot(train_loss, label=f'Train Run {i}')
    plt.plot(val_loss, label=f'Val Run {i}')

plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Train/Val Loss Across 3 Runs')
plt.legend()
plt.grid()
plt.tight_layout()
plt.show()
```




```
[ ]: # added this from a different file# to visualize the data distribution and
      ↪sample images (meant to be at the start of the notebook)

import pandas as pd
import matplotlib.pyplot as plt
from PIL import Image
import os

train_df = pd.read_csv('train_split.csv')


fig, axes = plt.subplots(1, 3, figsize=(15, 4))
targets = ['sog_speed', 'cog_speed', 'waterdepth']
titles = ['Speed over Ground (SOG)', 'Course over Ground (COG)', 'Water Depth_
      ↪(m)']

for ax, col, title in zip(axes, targets, titles):
    ax.hist(train_df[col].dropna(), bins=30)
    ax.set_title(f'Distribution of {title}')
    ax.set_xlabel(title)
    ax.set_ylabel('Count')

plt.tight_layout()
plt.show()

# Display the first frame of the first 9 sequences
fig, axes = plt.subplots(3, 3, figsize=(8, 8))
for i in range(9):
    row = train_df.iloc[i]
    img_path = os.path.join(
        'advanceaimethodsallfilestobeextracted22222222',
        row['subset'],
        row['filepath']
    )
    img = Image.open(img_path).convert('L').resize((128,128))
    ax = axes[i//3, i%3]
    ax.imshow(img, cmap='gray')
    ax.axis('off')
    ax.set_title(f"Seq #{i+1}")

plt.suptitle('Sample Sonar Image Frames (First of Each Sequence)')
plt.tight_layout()
plt.show()
```



Sample Sonar Image Frames (First of Each Sequence)

